

Notes on Probability Mass Functions

Probability Mass Functions

Note: I'm not going to refer to the sample space Ω in the main body of these notes.

In class we talked about random variables X , which we think of formally as numerical summaries of a random process. For example:

- You roll a die several times and want to summarize this process with X , the number of unique faces that we saw.
- You enter a raffle many times, and want to summarize this process with X , the amount of money that you won overall.
- You record the license plates of cars that drive by your apartment and summarize this process with X , the proportion of license plates that come from other states (not California).

Just as a random process can give you a different outcome each time you repeat it, the random variable X that summarizes that process can take on different values. We call the set of values that X can potentially take the *range* or *support* of X , and write it as \mathcal{X} . The *distribution* of X defines how likely X is to take any one of the values in its range \mathcal{X} . For the examples we consider in this class, where the range \mathcal{X} is finite (that is, X can only take finitely many values), we can represent this distribution using a *probability mass function*, or pmf for short¹.

A probability mass function (pmf) takes any value $k \in \mathcal{X}$ (that is, any value, which we will write as k , from the range \mathcal{X}), and returns the probability that X will take that value. We write this function using the notation $P(X = k)$.

Note that the symbol k here is just a placeholder variable for the *argument* of the pmf. We could equivalently use x or \heartsuit here, and write the same pmf as $P(X = x)$ or $P(X = \heartsuit)$. The important thing is that the pmf is a *function* of the placeholder variable.

Example: Binomial Random Variable

In class, we talked a binomial random variable X that we define as follows: Flip a coin n times. Each flip of the coin is independent, and has probability p of coming up heads. Define X as the number of heads that we see after flipping the coin n times.

In this case, the range of X is $\mathcal{X} = \{0, \dots, n\}$, since we can see as few as 0 heads or as many as n heads. The probability mass function for X is

$$P(X = k) = \binom{n}{k} p^k (1-p)^{n-k} \quad , \text{ or equivalently, } \quad P(X = \heartsuit) = \binom{n}{\heartsuit} p^{\heartsuit} (1-p)^{n-\heartsuit}$$

¹The range of \mathcal{X} of a random variable X is *not* the same as the sample space Ω . Ω is the *domain* of X ; X takes outcomes in the sample space Ω and summarizes them with a number contained in its range \mathcal{X} . In general, Ω is a set of full descriptions of ways that the random process could have turned out (e.g., sequences of heads and tails), but \mathcal{X} is a set of numbers.

for any $k \in \mathcal{X}$ (or any $\heartsuit \in \mathcal{X}$).

Using this function, you can answer any question of the form “What is the probability that we see m heads?” by plugging in m for k and evaluating the expression. For example, if n is greater than or equal to 10 and we ask, “What is the probability that we see 10 heads?”, we can plug in 10 for k and get our answer:

$$P(X = 10) = \binom{n}{10} p^{10} (1-p)^{n-10}.$$


Visualization

Probability mass functions are useful for visualizing the distributions of random variables. For each value $k \in \mathcal{X}$, we can plot the probability that X takes this value. Thus, the probability mass function tells you exactly how to draw a distribution plot. It can be intuitive to think that whenever we define a probability mass function, we are actually defining one of these plots.

Example: Binomial Distribution Plots

Using the pmf defined in the last example, if we supply a number of coin flips n and a probability of flipping heads p , we can plot the distribution of X , the number of heads that we would see if we flipped this coin n times.

For example, here are the distribution plots for three different settings of n and p :



These visualizations help us make observations like the following. When we run a coin flipping procedure like the one described:

- In all three cases, the most likely value that X (the number of heads) will take is np .
- It is very unlikely that we'll see X take a value in certain parts of the range \mathcal{X} . For example, in both cases where $n = 100$, it's very rare to see a number of heads that differs from the most likely value np by more than 10 heads.